a16z Big Ideas 2026: The Predictions Reshaping Growth, Infrastructure, and the Internet Itself

---

1. Source Info

Title: Big Ideas 2026 (Parts 1, 2, and 3)
Author: Andreessen Horowitz (a16z) Partners across Infrastructure, Growth, Bio + Health, Speedrun, American Dynamism, Apps, and Crypto teams
Date: March 2026

- **Links:**
- Part 1: https://a16z.com/newsletter/big-ideas-2026-part-1/
- Part 2: https://a16z.com/newsletter/big-ideas-2026-part-2/
- Part 3: https://a16z.com/newsletter/big-ideas-2026-part-3/

---

2. Executive Summary

Across three installments of their annual "Big Ideas" series, over 40 a16z partners lay out 22 predictions for 2026 spanning AI infrastructure, consumer applications, industrial reinvention, financial systems, and crypto. The throughline: AI is transitioning from a tool humans use to an actor that operates autonomously within and across systems. Infrastructure must be rebuilt for agent workloads (recursive, bursty, massively parallel). Products must be designed for two audiences simultaneously (humans and agents). Systems of record lose strategic primacy as AI collapses the intent-to-execution gap. And financial infrastructure must support machine-to-machine payments at protocol speed. For growth practitioners, the implications cut across every layer of the stack: the interfaces your users interact with, the metrics you optimize for, the channels you acquire through, and the monetization models you charge with are all being rearchitected around non-human actors as first-class participants.

3. Key Insights & Frameworks

3.1 Agent-Native Infrastructure: Rebuilding the Stack for Non-Human Users

Malika Aubakirova argues the entire infrastructure stack must be rebuilt for agent workloads. Today's infrastructure is optimized for human-speed interactions: predictable, low-concurrency, request-response patterns with 1:1 ratios of users to sessions.

Agent workloads are fundamentally different. They are recursive (an agent can spawn sub-agents that spawn their own sub-agents). They are bursty (an agent might sit idle for hours, then fire 10,000 API calls in 30 seconds). They demand massive parallelism (a single user might have dozens of agents running simultaneously). And they require coordination primitives that don't exist in most current stacks: routing, locking, state management, and policy enforcement across thousands of concurrent agent processes.

Aubakirova argues that cold starts must shrink to near-zero, latency variance must collapse, and concurrency limits must increase by orders of magnitude. The bottleneck becomes coordination, not compute. Gartner projects 40% of enterprise applications will feature task-specific AI agents by end of 2026, up from less than 5% in 2025.

3.2 The Death of the System of Record

Sarah Wang argues that CRM, ERP, and similar platforms are losing strategic primacy. When AI agents can read, write, and reason directly across operational data, the system of record becomes a background persistence layer, an infrastructure commodity rather than a strategic asset. The new strategic battleground is the "intelligent execution environment": the layer where AI agents interpret intent, orchestrate actions, and deliver outcomes.

Vendasta's research shows AI-native CRMs are closing the "execution gap," with organizations reporting 20-25% improvements in cycle times and measurable conversion rate increases. The CRM didn't go away; it just stopped being the thing that matters most.

3.3 Designing for Agents, Not Humans

Stephenie Zhang predicts web optimization is shifting from human consumption to agent consumption. Machine legibility replaces visual hierarchy: structured data, clean APIs, semantic markup, and fast response times. AI agents typically have timeouts of one to five seconds for retrieving information, so speed directly affects visibility.

SEO leaders note the industry is splitting into two distinct strategic problems. Traditional SEO focuses on humans who browse, compare, and buy. AI search optimization focuses on supplying information so agents can find, trust, and use it without a human ever visiting the site. If your product data, pricing, and availability are not machine-readable in real time, AI agents will skip you.

This creates the dual-audience challenge: your website needs to convert human visitors through compelling design AND serve agent visitors through structured data and fast APIs. These are not conflicting goals, but they require different optimization strategies running in parallel.

3.4 Beyond Screen Time: New KPIs for the Agent Era

Santiago Rodriguez argues that engagement metrics (DAU, time spent, sessions per week) assume value correlates with attention. AI applications increasingly deliver value without consuming screen time. A doctor using an AI diagnostic tool gets the answer in seconds. A developer using Claude Code has agents running in the background.

Companies measuring ROI by outcomes (problems resolved, productivity gained) rather than engagement metrics will outpace competitors still optimizing for attention. The best AI products will be the ones users spend the least time using, because the AI handles everything so efficiently that extended interaction is unnecessary.

3.5 ChatGPT as the New App Store

Anish Acharya frames ChatGPT (approximately 900 million weekly active users) as a consumer distribution platform comparable to the early iOS App Store. Apps built on OpenAI's SDK (based on Model Context Protocol) surface automatically within conversations

based on user intent. No download, no separate app launch, no traditional customer acquisition cost.

Three distribution mechanics change: (1) Discovery shifts from search-based to intent-based; users describe a problem and apps appear if they solve it. (2) The conversion funnel collapses; no download step, no onboarding flow, no account creation barrier. Users go from intent to value in a single conversational turn. (3) Retention becomes relevance-based rather than habit-based; users don't "return to your app" but return to ChatGPT and your app resurfaces when relevant.

3.6 Prompt-Free Applications: AI Becomes Invisible

Marc Andrusko argues the prompt-based AI paradigm (type instructions into a text box) is transitional. In 2026, AI increasingly intervenes proactively through observation rather than waiting for prompts, integrated into IDEs, CRMs, and design tools, surfacing suggestions and taking actions based on context.

Jakob Nielsen's 2026 UX predictions align: the chat box is no longer the default solution. Designers are treating AI as an invisible layer powering the entire application. The winners are vertical AI platforms wrapping commoditized models in highly specific, defensible workflows.

For growth teams, prompt-free AI changes the activation equation. In prompt-based products, the biggest barrier is the "blank prompt problem": users don't know what to ask. Prompt-free products eliminate this by proactively delivering value. Activation rates should be dramatically higher, but the product must be smart enough to deliver relevant interventions without being asked. Getting this wrong (surfacing irrelevant or annoying suggestions) destroys trust faster than a blank prompt box does.

3.7 The Year of Personalization

Joshua Lu predicts 2026 is "the year of me": products stop being mass-produced and start being individually tailored. AI tutors adapt to learning styles, health plans customize to biomarkers, media feeds are curated by AI that understands individual preferences at a deep level.

This is not the personalization of 2020 (showing different banner ads to different segments). This is hyper-personalization at the individual product experience level. The global hyper-personalization market was valued at approximately $21.2 billion in 2024, with projections reaching $68 billion by 2031. Research shows 81% of consumers prefer companies offering personalized experiences.

For growth teams, this changes the experimentation model. Instead of A/B tests (Treatment A vs. Treatment B for 50/50 splits), you can run N-of-1 experiments where every user gets a uniquely tailored experience. The challenge shifts from "which variant wins?" to "how do we build the personalization engine that generates optimal variants for each individual?"

The privacy dimension is critical. The most successful personalization in 2026 treats privacy as a premium feature, moving from third-party tracking to explicit data collection where users voluntarily share preferences because they see direct value in return.

3.8 Voice Agents: From Point Solutions to Full Workflow Management

Olivia Moore predicts voice agents will evolve from handling isolated tasks to managing entire customer workflows: appointments, bookings, intake forms, follow-up sequences, and full customer lifecycle management.

The voice AI market is projected to grow from $2.4 billion in 2024 to $47.5 billion by 2034. One in ten customer service interactions will be fully automated by agentic voice AI by 2026. Healthcare alone is projected to save $150 billion annually through voice AI.

Teams are running hybrid human/AI models rather than pursuing full automation. Growth teams need to optimize the handoff between AI and human agents, measuring where AI resolution is sufficient and where human intervention drives better outcomes.

3.9 Multi-Agent Orchestration in the Enterprise

Seema Amble predicts Fortune 500 companies will deploy coordinated multi-agent systems managing complex workflows. IBM's research shows multi-agent orchestration reduces handoffs by 45% and increases decision speed by 3x. The analogy: orchestration layers for agents are comparable to what Kubernetes did for container management.

Amble predicts the emergence of new roles: "AI workflow designers" who architect agent systems, and "agent supervisors" who monitor and tune agent performance.

3.10 The American Industrial Renaissance

The American Dynamism team predicts an AI-native industrial base. David Ulevitch envisions manufacturing and infrastructure companies starting with simulation and AI-driven operations. Erin Price-Wright sees the factory itself becoming the product.

The economics are real. The U.S. faces a projected labor deficit of 425,000 workers in construction and advanced manufacturing. Electricity costs are rising 18.3% year-over-year in key manufacturing corridors. Global smart manufacturing adoption has reached 47%, up 12 percentage points from the prior year. ABB and NVIDIA have closed the simulation-to-reality gap in industrial robotics, achieving 99% correlation between simulated and real-world robot behavior.

For growth practitioners in B2B SaaS, this represents an enormous market opportunity. The industrial sector is traditionally underserved by modern software, and companies are being forced to modernize by labor shortages and energy costs.

Zabie Elmgren's "physical observability" prediction (networked cameras and sensors enabling real-time infrastructure monitoring) creates data volumes that dwarf traditional SaaS metrics. Growth engineering principles (instrumentation, experimentation, optimization loops) apply directly to physical infrastructure when sensors make the physical world measurable in real time.

3.11 The Electro-Industrial Stack

Ryan McEntush introduces the "electro-industrial stack": the integration of electrification hardware (batteries, motors, power electronics) with software control layers. Every battery, motor, and power system increasingly has a software control plane. The companies building this control software are essentially building growth-engineered products for physical infrastructure: optimizing performance through data, running experiments on control algorithms, measuring outcomes in real-world metrics (energy efficiency, uptime, cost per unit).

3.12 Multimodal Data and the AI-Native Data Stack

Jennifer Li predicts enterprises will invest heavily in extracting structure from unstructured data: PDFs, videos, logs, emails, images. Jason Cui adds that the data stack itself is consolidating, with opportunities in vector databases, agentic BI tools, and AI agents solving "the context problem."

For growth engineers, data stack consolidation is both opportunity and warning. Opportunity: fewer, more integrated tools lower experimentation costs and speed insight. Warning: elaborate integrations between specialized tools may be rendered obsolete. Growth teams sit on enormous unstructured data (user feedback, support tickets, sales call transcripts). AI platforms that extract actionable structure from this data give growth teams qualitative insight at quantitative scale.

3.13 Video as Inhabitable Space

Yoko Li predicts video stops being something we passively watch and starts being a place we step into. Video models that sustain characters and physics over time enable interactive environments. Jonathan Lai references world-generating AI (Marble, Genie 3) creating interactive 3D environments: "generative Minecraft" where worlds are created and inhabited simultaneously.

For growth teams, interactive media represents a new product category that does not fit existing engagement frameworks. If a user spends 30 minutes in a generative 3D environment, is that video consumption, gaming, or product usage? Growth teams building here will need to invent new measurement frameworks.

3.14 Healthy MAUs: The Wellness Consumer Segment

Julie Yoo predicts AI-enabled preventive care creates a new market segment: health-conscious consumers who want recurring health services even when not sick. This is a TAM expansion argument; the addressable market shifts from "people who are currently sick" to "people who want to stay healthy."

The "Healthy MAU" concept is transferable beyond healthcare. It represents a pattern: using

AI to shift a product's value proposition from reactive (solving a problem when it occurs) to proactive (preventing the problem or enhancing ongoing performance). Any product making this shift dramatically expands its addressable market and changes retention dynamics from episodic to continuous.

3.15 The First AI-Native University

Emily Bennett predicts universities built around AI systems: adaptive courses adjusting to individual learning pace, real-time learning paths, and professors as "architects of learning." If AI-native education produces graduates expecting adaptive, personalized learning, they will also expect products that adapt to them. Growth teams building adaptive onboarding (adjusting the experience based on user behavior and knowledge level) will have an activation advantage over static tutorials.

3.16 Know Your Agent: Identity for Non-Human Actors

Sean Neville introduces "Know Your Agent" (KYA). In financial services, there are already 96 non-human identities for every human employee. KYA proposes cryptographically signed credentials linking agents to their principals (who deployed them), constraints (what they can do), and liability (who is responsible for mistakes).

3.17 Programmable Money and Machine-to-Machine Payments

Christian Crowley predicts programmable money through primitives like HTTP 402: agents paying other agents for data, compute, and API access without human intermediation.

Real infrastructure supports this. Coinbase's x402 protocol embeds stablecoin payments directly into HTTP requests. Stripe-backed Tempo launched a Machine Payments Protocol for AI agent transactions. Stablecoins have reached $46 trillion in transaction volume (20x PayPal, 3x Visa).

3.18 The Invisible Tax on the Open Web

Elizabeth Harkavy warns about an "invisible tax" AI agents impose on the ad-supported web. When agents browse on behalf of users, they bypass ads that fund content creation. Publishers get traffic but no revenue.

Harkavy proposes blockchain-based nanopayment systems for real-time compensation. Google has launched Offerwall to help publishers beyond ads. OpenAI is introducing advertising to ChatGPT (approximately $60 per 1,000 impressions) with revenue sharing for content producers.

3.19 Privacy as Crypto's Competitive Moat

Ali Yahya argues privacy will become the strongest competitive moat in crypto. Bridging secrets between blockchains is harder than moving public tokens. If private data lives on a specific chain, switching requires exposing it during migration, creating "chain lock-in" stronger than feature-based differentiation.

This is a growth insight disguised as a crypto prediction. Any product storing sensitive user data creates lock-in because migration requires re-establishing trust with a new provider. Zero-knowledge proofs now enable compliance-friendly privacy where users prove they meet regulatory requirements without revealing underlying data. For growth teams in regulated industries, ZK-based verification could dramatically simplify compliance-heavy onboarding flows.

3.20 Prediction Markets Go Mainstream

Andrew Hall predicts prediction markets will expand dramatically. Kalshi reached $22 billion valuation; Polymarket hit $20 billion. Over 30% of wallets on Polymarket use AI agents. Olas's Polystrat agent executed over 4,200 trades in a single month, achieving returns as high as 376% on individual trades.

For growth teams, internal prediction markets can complement A/B testing: before running an experiment, let team members trade on the expected outcome. This calibrates priors, surfaces hidden knowledge, and creates accountability for predictions. Companies like Google

have used internal prediction markets for years; AI agents make it possible to scale this approach.

3.21 Staked Media and Credibility Infrastructure

Robert Hackett introduces "staked media": journalists and analysts lock tokens to prove financial alignment with their claims. Credibility becomes auditable through transparent commitment rather than reputation alone.

The underlying principle is transferable: how can you make your claims verifiable? If your marketing says "our product increases conversion by 30%," can you back that with an auditable commitment? In a world of abundant AI-generated content, credibility becomes a scarce and valuable asset.

3.22 Forward-Deployed Motions and Greenfield Strategy

Two go-to-market predictions from the Apps team. Joe Schmidt argues for "forward-deployed motions": startups discovering opportunities through direct engagement outside Silicon Valley, working with traditional industries. James da Costa advocates a "greenfield strategy": targeting newly formed companies rather than displacing incumbents.

The greenfield strategy follows the Stripe and Deel playbook: rather than competing for existing customers, target companies being formed right now with no legacy systems. Switching cost for a company that hasn't adopted any tool is zero. The forward-deployed motion is the inverse: going where software is underserved and building relationships through direct engagement. Both recognize that the highest-leverage distribution is finding customers who need what you offer but have not yet been served.

---

4. Tactical Playbook

4.1 Audit Your Infrastructure for Agent Readiness (see 3.1)

Review APIs, analytics pipelines, and experimentation platforms. Ask: can they handle 100x current volume with 10x lower latency? If an AI agent were running your experiment pipeline programmatically, what would break first? Fix that thing.

4.2 Build a Machine-Readable Layer for Your Product (see 3.3)

Add structured data, clean APIs, and semantic markup to web properties. Create an API endpoint surfacing your product's key capabilities, pricing, and availability in a format agents can consume. This is "agent SEO" and it should be on your roadmap now.

4.3 Prototype a ChatGPT App (see 3.5)

Build a minimal ChatGPT app using OpenAI's SDK. Even a simple integration puts you on a distribution layer with massive reach. Measure intent-based discovery and conversion rates compared to traditional acquisition channels.

4.4 Shift Your Metric Framework from Engagement to Outcomes (see 3.4)

Identify the primary outcome your product delivers. Build dashboards tracking "value delivered per unit of user attention" rather than just "time spent." If your product can deliver value in 30 seconds instead of 30 minutes, your metrics should celebrate that efficiency.

4.5 Experiment with Voice Agent Acquisition (see 3.8)

Set up a voice agent to handle one specific acquisition workflow: inbound lead qualification, demo scheduling, or trial activation. Measure conversion rates and cost per acquisition against your current approach. The economics make this a low-risk experiment with potentially high payoff.

4.6 Design Your Next Feature as Prompt-Free (see 3.6)

Instead of adding a chat interface or prompt box, build a feature that observes user behavior and proactively surfaces relevant information. Measure whether proactive intervention achieves higher adoption than a feature requiring users to know to ask.

4.7 Map Your Agent Authentication Strategy (see 3.16)

If your product exposes APIs, prototype a KYA-style authentication system identifying whether the caller is a human, an agent, or a bot. This allows differentiated experiences and pricing by actor type, positioning you for agent traffic exceeding human traffic.

4.8 Target Greenfield Customers (see 3.22)

Identify companies formed in the last 6 months matching your ideal customer profile. Build an outreach workflow targeting these nascent companies before they adopt a competitor. Switching cost for a company that hasn't adopted any tool is zero.

4.9 Build a Qualitative Data Pipeline (see 3.12)

Use AI to extract structured insights from unstructured data: support tickets, user feedback, sales call transcripts, product reviews. Feed these insights into your experiment pipeline. The gap between qualitative understanding and quantitative experimentation should be closing.

4.10 Run an Internal Prediction Market (see 3.20)

Before your next major experiment, let team members predict the outcome. Compare predictions against actual results. This calibrates judgment, surfaces hidden assumptions, and creates a culture of accountable prediction that improves experiment design over time.

---

5. Contrarian / Non-Obvious Takes

5.1 Less Screen Time Means More Value

The entire consumer tech industry was built on the assumption that engagement (measured by screen time) equals value. Rodriguez's prediction directly inverts this. But the deeper implication goes beyond metrics: growth teams optimizing for reduced screen time must also rethink monetization. Attention-based business models (ad-supported, engagement-driven upsells) break when the product's job is to minimize user time. The companies that figure out outcome-based pricing for low-attention products will define the next generation of growth models.

5.2 Privacy Creates Stronger Lock-In Than Features

Most product teams treat privacy as a cost center. But Yahya's insight goes further than defensibility: privacy-as-lock-in compounds over time. Every additional piece of sensitive data a user entrusts to your platform increases the switching cost without you building anything new. This suggests a counterintuitive growth strategy: actively encourage users to store their most sensitive data with you, not for monetization, but to deepen the structural moat. The products handling your most sensitive data are the hardest to leave.

5.3 The Factory Is the Product

Price-Wright's claim is counterintuitive for an industry that thinks of software as the product and hardware as a commodity. The deeper implication: the defensible asset in the industrial renaissance is not IP or algorithms but the physical-digital integration itself. A factory with 99% sim-to-real correlation in its robotic systems is a competitive moat that cannot be replicated with software alone. This challenges the "software eats everything" thesis by arguing that some of the most valuable systems are ones where software and atoms are inseparable.

5.4 Your Next Competitor Is an Agent, Not a Company

Across multiple predictions (3.5, 3.9, 3.16), a pattern emerges that none of the individual partners state explicitly: the competitive landscape is shifting from company-versus-company to product-versus-agent. An agent that can orchestrate multiple APIs to solve a user's problem competes with your all-in-one product even if no single company built that agent to compete with you. The unbundling threat comes not from startups but from AI agents that compose solutions dynamically from available APIs.

5.5 Staking Your Reputation with Tokens

Hackett's staked media concept inverts the current credibility model from reputation-based (track record and social proof) to commitment-based (financial skin in the game). The uncomfortable implication: if verifiable commitment becomes the credibility standard, the value of thought leadership, conference talks, and newsletter authority diminishes relative to people willing to put money behind their predictions. This restructures the incentives for growth marketing content from "be interesting" to "be correct and prove it."

---

6. What to Revisit

Part 1, Malika Aubakirova's agent-native infrastructure prediction: Read carefully if building or managing infrastructure that AI agents will interact with. The specific technical requirements (near-zero cold starts, collapsed latency variance, massive concurrency) are actionable design constraints.

Part 1, Sarah Wang on systems of record: Revisit when evaluating your analytics and CRM stack. Consider whether your current system of record is positioned to become infrastructure or strategy.

Part 1, Stephenie Zhang on designing for agents: Read alongside your SEO and web optimization strategy. Assess whether your web properties are machine-readable.

Part 2, Anish Acharya on ChatGPT as app store: Read before your next channel strategy meeting. Evaluate whether building a ChatGPT app should be on your roadmap.

Part 2, Olivia Moore on voice agents: Revisit when planning customer lifecycle automation. Evaluate where voice agents could replace or augment human-mediated workflows.

Part 2, Seema Amble on multi-agent orchestration: Read when designing complex workflows. Consider whether agent orchestration could automate your growth experiment pipeline.

Part 3, Sean Neville on KYA: Revisit when designing API authentication. Consider how your product should identify and serve agent users differently from human users.

Part 3, Elizabeth Harkavy on the invisible tax: Read when evaluating content marketing ROI. Assess how much of your content is being consumed by agents that never convert.

Part 3, Ali Yahya on privacy as moat: Revisit when evaluating retention and switching costs. Consider whether your handling of sensitive data creates structural lock-in.

Part 3, Andrew Hall on prediction markets: Read when evaluating experiment design. Consider whether internal prediction markets could improve your team's forecasting accuracy.

---

7. Sources & Further Reading

1. **a16z Big Ideas 2026, Part 1**

https://a16z.com/newsletter/big-ideas-2026-part-1/ Infrastructure, Growth, Bio + Health, and Speedrun team predictions. Covers agent-native infrastructure, systems of record, video as inhabitable space, and personalization.

2. **a16z Big Ideas 2026, Part 2**

https://a16z.com/newsletter/big-ideas-2026-part-2/ American Dynamism and Apps team predictions. Covers industrial AI, factory renaissance, ChatGPT as app store, voice agents, and multi-agent orchestration.

3. **a16z Big Ideas 2026, Part 3**

https://a16z.com/newsletter/big-ideas-2026-part-3/ Crypto team predictions. Covers KYA, programmable money, privacy as moat, prediction markets, staked media, and regulation.

4. **Gartner: 40% of Enterprise Apps to Feature AI Agents by 2026**

https://www.gartner.com/en/newsroom/press-releases/2025-08-26-gartner-predicts-40-percent-of-enterprise-apps Industry benchmark data validating the agent infrastructure predictions.

5. **a16z: The Top 100 Gen AI Consumer Apps, 6th Edition**

https://a16z.com/100-gen-ai-apps-6/ Contextualizes the ChatGPT app store within the broader consumer AI landscape and distribution dynamics.

6. **TechCrunch: ChatGPT Launches an App Store**

https://techcrunch.com/2025/12/18/chatgpt-launches-an-app-store-lets-developers-know-its-open-for-business/ Details on the technical architecture and distribution mechanics of ChatGPT's app platform.

7. **Search Engine Land: The Future of AI Search, 6 SEO Leaders Predict 2026**

https://searchengineland.com/ai-search-visibility-seo-predictions-2026-468042 Detailed analysis of how web optimization splits into human-focused and agent-focused strategies.

8. **Vendasta: AI-Native CRM and the End of the Execution Gap**

https://www.vendasta.com/blog/ai-native-crm/ Data on how AI-native CRM approaches are closing the gap between data capture and autonomous action.

9. **CoinDesk: AI Agents Are Quietly Rewriting Prediction Market Trading**

https://www.coindesk.com/tech/2026/03/15/ai-agents-are-quietly-rewriting-prediction-market-trading Details on AI agent participation in prediction markets, including Polystrat's performance data.

10. **CoinDesk: Visa, Coinbase Building Different Internets for AI Agent Payments**

https://www.coindesk.com/tech/2026/03/15/visa-is-ready-for-ai-agents-so-is-coinbase-they-re-building-very-differe Infrastructure details on how stablecoins and traditional payment rails are adapting for machine-to-machine payments.

11. **Manufacturing Dive: Physical AI and Automation Trends for 2026**

https://www.manufacturingdive.com/news/physical-ai-craze-2026-automation-trends-to-watch/810860/ Market data and case studies on the American industrial renaissance and factory automation.

12. **SignalFire: Seven Frameworks for AI-Enabled Vertical SaaS**

https://www.signalfire.com/blog/frameworks-for-ai-vertical-saas Growth frameworks for building

and scaling vertical AI companies, directly applicable to the multiplayer AI prediction.

13. **O'Reilly: Signals for 2026**

https://www.oreilly.com/radar/signals-for-2026/ Broader technology landscape context for the agent infrastructure and AI-native development predictions.

14. **Jakob Nielsen: 18 Predictions for 2026**

https://jakobnielsenphd.substack.com/p/2026-predictions UX perspective on prompt-free AI and the shift from chat interfaces to invisible AI integration.

15. **Fortune: Prediction Market Boom Spurs New VC Fund**

https://fortune.com/2026/03/23/kalshi-polymarket-5cc-capital-prediction-market-fund-raise/ Context on the scale and institutional validation of prediction markets as a growing asset class.